

G-29_B_report

April 17, 2024

```
[1]: !pip install rouge_score
```

```
Collecting rouge_score
  Downloading rouge_score-0.1.2.tar.gz (17 kB)
  Preparing metadata (setup.py) ... done
Requirement already satisfied: absl-py in /usr/local/lib/python3.10/dist-packages (from rouge_score) (1.4.0)
Requirement already satisfied: nltk in /usr/local/lib/python3.10/dist-packages (from rouge_score) (3.8.1)
Requirement already satisfied: numpy in /usr/local/lib/python3.10/dist-packages (from rouge_score) (1.25.2)
Requirement already satisfied: six>=1.14.0 in /usr/local/lib/python3.10/dist-packages (from rouge_score) (1.16.0)
Requirement already satisfied: click in /usr/local/lib/python3.10/dist-packages (from nltk->rouge_score) (8.1.7)
Requirement already satisfied: joblib in /usr/local/lib/python3.10/dist-packages (from nltk->rouge_score) (1.4.0)
Requirement already satisfied: regex>=2021.8.3 in /usr/local/lib/python3.10/dist-packages (from nltk->rouge_score) (2023.12.25)
Requirement already satisfied: tqdm in /usr/local/lib/python3.10/dist-packages (from nltk->rouge_score) (4.66.2)
Building wheels for collected packages: rouge_score
  Building wheel for rouge_score (setup.py) ... done
  Created wheel for rouge_score: filename=rouge_score-0.1.2-py3-none-any.whl size=24933 sha256=a1e916bae74c0725f0cd28078e6e23029d42afabeab43092fab066454918c208
    Stored in directory: /root/.cache/pip/wheels/5f/dd/89/461065a73be61a532ff8599a28e9beef17985c9e9c31e541b4
Successfully built rouge_score
Installing collected packages: rouge_score
Successfully installed rouge_score-0.1.2
```

```
[2]: !pip install datasets
```

```
Collecting datasets
  Downloading datasets-2.18.0-py3-none-any.whl (510 kB)
```

510.5/510.5

kB 5.9 MB/s eta 0:00:00

Requirement already satisfied: filelock in /usr/local/lib/python3.10/dist-packages (from datasets) (3.13.4)

Requirement already satisfied: numpy>=1.17 in /usr/local/lib/python3.10/dist-packages (from datasets) (1.25.2)

Requirement already satisfied: pyarrow>=12.0.0 in /usr/local/lib/python3.10/dist-packages (from datasets) (14.0.2)

Requirement already satisfied: pyarrow-hotfix in /usr/local/lib/python3.10/dist-packages (from datasets) (0.6)

Collecting dill<0.3.9,>=0.3.0 (from datasets)

Downloading dill-0.3.8-py3-none-any.whl (116 kB)

116.3/116.3

kB 6.1 MB/s eta 0:00:00

Requirement already satisfied: pandas in /usr/local/lib/python3.10/dist-packages (from datasets) (2.0.3)

Requirement already satisfied: requests>=2.19.0 in /usr/local/lib/python3.10/dist-packages (from datasets) (2.31.0)

Requirement already satisfied: tqdm>=4.62.1 in /usr/local/lib/python3.10/dist-packages (from datasets) (4.66.2)

Collecting xxhash (from datasets)

Downloading

xxhash-3.4.1-cp310-cp310-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (194 kB)

194.1/194.1

kB 10.2 MB/s eta 0:00:00

Collecting multiprocessing (from datasets)

Downloading multiprocessing-0.70.16-py310-none-any.whl (134 kB)

134.8/134.8

kB 10.6 MB/s eta 0:00:00

Requirement already satisfied: fsspec[http]<=2024.2.0,>=2023.1.0 in /usr/local/lib/python3.10/dist-packages (from datasets) (2023.6.0)

Requirement already satisfied: aiohttp in /usr/local/lib/python3.10/dist-packages (from datasets) (3.9.3)

Requirement already satisfied: huggingface-hub>=0.19.4 in /usr/local/lib/python3.10/dist-packages (from datasets) (0.20.3)

Requirement already satisfied: packaging in /usr/local/lib/python3.10/dist-packages (from datasets) (24.0)

Requirement already satisfied: pyyaml>=5.1 in /usr/local/lib/python3.10/dist-packages (from datasets) (6.0.1)

Requirement already satisfied: aiosignal>=1.1.2 in /usr/local/lib/python3.10/dist-packages (from aiohttp->datasets) (1.3.1)

Requirement already satisfied: attrs>=17.3.0 in /usr/local/lib/python3.10/dist-packages (from aiohttp->datasets) (23.2.0)

Requirement already satisfied: frozenlist>=1.1.1 in /usr/local/lib/python3.10/dist-packages (from aiohttp->datasets) (1.4.1)

Requirement already satisfied: multidict<7.0,>=4.5 in /usr/local/lib/python3.10/dist-packages (from aiohttp->datasets) (6.0.5)
 Requirement already satisfied: yarl<2.0,>=1.0 in /usr/local/lib/python3.10/dist-packages (from aiohttp->datasets) (1.9.4)
 Requirement already satisfied: async-timeout<5.0,>=4.0 in /usr/local/lib/python3.10/dist-packages (from aiohttp->datasets) (4.0.3)
 Requirement already satisfied: typing-extensions>=3.7.4.3 in /usr/local/lib/python3.10/dist-packages (from huggingface-hub>=0.19.4->datasets) (4.11.0)
 Requirement already satisfied: charset-normalizer<4,>=2 in /usr/local/lib/python3.10/dist-packages (from requests>=2.19.0->datasets) (3.3.2)
 Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.10/dist-packages (from requests>=2.19.0->datasets) (3.6)
 Requirement already satisfied: urllib3<3,>=1.21.1 in /usr/local/lib/python3.10/dist-packages (from requests>=2.19.0->datasets) (2.0.7)
 Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.10/dist-packages (from requests>=2.19.0->datasets) (2024.2.2)
 Requirement already satisfied: python-dateutil>=2.8.2 in /usr/local/lib/python3.10/dist-packages (from pandas->datasets) (2.8.2)
 Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.10/dist-packages (from pandas->datasets) (2023.4)
 Requirement already satisfied: tzdata>=2022.1 in /usr/local/lib/python3.10/dist-packages (from pandas->datasets) (2024.1)
 Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.10/dist-packages (from python-dateutil>=2.8.2->pandas->datasets) (1.16.0)
 Installing collected packages: xxhash, dill, multiprocessing, datasets
 Successfully installed datasets-2.18.0 dill-0.3.8 multiprocessing-0.70.16 xxhash-3.4.1

```
[3]: !pip install transformers[torch]
```

Requirement already satisfied: transformers[torch] in /usr/local/lib/python3.10/dist-packages (4.38.2)
 Requirement already satisfied: filelock in /usr/local/lib/python3.10/dist-packages (from transformers[torch]) (3.13.4)
 Requirement already satisfied: huggingface-hub<1.0,>=0.19.3 in /usr/local/lib/python3.10/dist-packages (from transformers[torch]) (0.20.3)
 Requirement already satisfied: numpy>=1.17 in /usr/local/lib/python3.10/dist-packages (from transformers[torch]) (1.25.2)
 Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.10/dist-packages (from transformers[torch]) (24.0)
 Requirement already satisfied: pyyaml>=5.1 in /usr/local/lib/python3.10/dist-packages (from transformers[torch]) (6.0.1)
 Requirement already satisfied: regex!=2019.12.17 in /usr/local/lib/python3.10/dist-packages (from transformers[torch]) (2023.12.25)

Requirement already satisfied: requests in /usr/local/lib/python3.10/dist-packages (from transformers[torch]) (2.31.0)
Requirement already satisfied: tokenizers<0.19,>=0.14 in /usr/local/lib/python3.10/dist-packages (from transformers[torch]) (0.15.2)
Requirement already satisfied: safetensors>=0.4.1 in /usr/local/lib/python3.10/dist-packages (from transformers[torch]) (0.4.2)
Requirement already satisfied: tqdm>=4.27 in /usr/local/lib/python3.10/dist-packages (from transformers[torch]) (4.66.2)
Requirement already satisfied: torch in /usr/local/lib/python3.10/dist-packages (from transformers[torch]) (2.2.1+cu121)
Collecting accelerate>=0.21.0 (from transformers[torch])
 Downloading accelerate-0.29.2-py3-none-any.whl (297 kB)
297.4/297.4

kB 6.4 MB/s eta 0:00:00

Requirement already satisfied: psutil in /usr/local/lib/python3.10/dist-packages (from accelerate>=0.21.0->transformers[torch]) (5.9.5)
Requirement already satisfied: fsspec>=2023.5.0 in /usr/local/lib/python3.10/dist-packages (from huggingface-hub<1.0,>=0.19.3->transformers[torch]) (2023.6.0)
Requirement already satisfied: typing-extensions>=3.7.4.3 in /usr/local/lib/python3.10/dist-packages (from huggingface-hub<1.0,>=0.19.3->transformers[torch]) (4.11.0)
Requirement already satisfied: sympy in /usr/local/lib/python3.10/dist-packages (from torch->transformers[torch]) (1.12)
Requirement already satisfied: networkx in /usr/local/lib/python3.10/dist-packages (from torch->transformers[torch]) (3.3)
Requirement already satisfied: Jinja2 in /usr/local/lib/python3.10/dist-packages (from torch->transformers[torch]) (3.1.3)
Collecting nvidia-cuda-nvrtc-cu12==12.1.105 (from torch->transformers[torch])
 Using cached nvidia_cuda_nvrtc_cu12-12.1.105-py3-none-manylinux1_x86_64.whl (23.7 MB)
Collecting nvidia-cuda-runtime-cu12==12.1.105 (from torch->transformers[torch])
 Using cached nvidia_cuda_runtime_cu12-12.1.105-py3-none-manylinux1_x86_64.whl (823 kB)
Collecting nvidia-cuda-cupti-cu12==12.1.105 (from torch->transformers[torch])
 Using cached nvidia_cuda_cupti_cu12-12.1.105-py3-none-manylinux1_x86_64.whl (14.1 MB)
Collecting nvidia-cudnn-cu12==8.9.2.26 (from torch->transformers[torch])
 Using cached nvidia_cudnn_cu12-8.9.2.26-py3-none-manylinux1_x86_64.whl (731.7 MB)
Collecting nvidia-cublas-cu12==12.1.3.1 (from torch->transformers[torch])
 Using cached nvidia_cublas_cu12-12.1.3.1-py3-none-manylinux1_x86_64.whl (410.6 MB)
Collecting nvidia-cufft-cu12==11.0.2.54 (from torch->transformers[torch])
 Using cached nvidia_cufft_cu12-11.0.2.54-py3-none-manylinux1_x86_64.whl (121.6 MB)
Collecting nvidia-curand-cu12==10.3.2.106 (from torch->transformers[torch])

```

Using cached nvidia_curand_cu12-10.3.2.106-py3-none-manylinux1_x86_64.whl
(56.5 MB)
Collecting nvidia-cusolver-cu12==11.4.5.107 (from torch->transformers[torch])
Using cached nvidia_cusolver_cu12-11.4.5.107-py3-none-manylinux1_x86_64.whl
(124.2 MB)
Collecting nvidia-cuspars-cu12==12.1.0.106 (from torch->transformers[torch])
Using cached nvidia_cuspars-cu12-12.1.0.106-py3-none-manylinux1_x86_64.whl
(196.0 MB)
Collecting nvidia-nccl-cu12==2.19.3 (from torch->transformers[torch])
Using cached nvidia_nccl_cu12-2.19.3-py3-none-manylinux1_x86_64.whl (166.0 MB)
Collecting nvidia-nvtx-cu12==12.1.105 (from torch->transformers[torch])
Using cached nvidia_nvtx_cu12-12.1.105-py3-none-manylinux1_x86_64.whl (99 kB)
Requirement already satisfied: triton==2.2.0 in /usr/local/lib/python3.10/dist-
packages (from torch->transformers[torch]) (2.2.0)
Collecting nvidia-nvjitlink-cu12 (from nvidia-cusolver-
cu12==11.4.5.107->torch->transformers[torch])
Using cached nvidia_nvjitlink_cu12-12.4.127-py3-none-manylinux2014_x86_64.whl
(21.1 MB)
Requirement already satisfied: charset-normalizer<4,>=2 in
/usr/local/lib/python3.10/dist-packages (from requests->transformers[torch])
(3.3.2)
Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.10/dist-
packages (from requests->transformers[torch]) (3.6)
Requirement already satisfied: urllib3<3,>=1.21.1 in
/usr/local/lib/python3.10/dist-packages (from requests->transformers[torch])
(2.0.7)
Requirement already satisfied: certifi>=2017.4.17 in
/usr/local/lib/python3.10/dist-packages (from requests->transformers[torch])
(2024.2.2)
Requirement already satisfied: MarkupSafe>=2.0 in
/usr/local/lib/python3.10/dist-packages (from
jinja2->torch->transformers[torch]) (2.1.5)
Requirement already satisfied: mpmath>=0.19 in /usr/local/lib/python3.10/dist-
packages (from sympy->torch->transformers[torch]) (1.3.0)
Installing collected packages: nvidia-nvtx-cu12, nvidia-nvjitlink-cu12, nvidia-
nccl-cu12, nvidia-curand-cu12, nvidia-cufft-cu12, nvidia-cuda-runtime-cu12,
nvidia-cuda-nvrtc-cu12, nvidia-cuda-cupti-cu12, nvidia-cublas-cu12, nvidia-
cuspars-cu12, nvidia-cudnn-cu12, nvidia-cusolver-cu12, accelerate
Successfully installed accelerate-0.29.2 nvidia-cublas-cu12-12.1.3.1 nvidia-
cuda-cupti-cu12-12.1.105 nvidia-cuda-nvrtc-cu12-12.1.105 nvidia-cuda-runtime-
cu12-12.1.105 nvidia-cudnn-cu12-8.9.2.26 nvidia-cufft-cu12-11.0.2.54 nvidia-
curand-cu12-10.3.2.106 nvidia-cusolver-cu12-11.4.5.107 nvidia-cuspars-
cu12-12.1.0.106 nvidia-nccl-cu12-2.19.3 nvidia-nvjitlink-cu12-12.4.127 nvidia-
nvtx-cu12-12.1.105

```

```
[4]: !pip install accelerate -U
```

```
Requirement already satisfied: accelerate in /usr/local/lib/python3.10/dist-
```

packages (0.29.2)
 Requirement already satisfied: numpy>=1.17 in /usr/local/lib/python3.10/dist-packages (from accelerate) (1.25.2)
 Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.10/dist-packages (from accelerate) (24.0)
 Requirement already satisfied: psutil in /usr/local/lib/python3.10/dist-packages (from accelerate) (5.9.5)
 Requirement already satisfied: pyyaml in /usr/local/lib/python3.10/dist-packages (from accelerate) (6.0.1)
 Requirement already satisfied: torch>=1.10.0 in /usr/local/lib/python3.10/dist-packages (from accelerate) (2.2.1+cu121)
 Requirement already satisfied: huggingface-hub in /usr/local/lib/python3.10/dist-packages (from accelerate) (0.20.3)
 Requirement already satisfied: safetensors>=0.3.1 in /usr/local/lib/python3.10/dist-packages (from accelerate) (0.4.2)
 Requirement already satisfied: filelock in /usr/local/lib/python3.10/dist-packages (from torch>=1.10.0->accelerate) (3.13.4)
 Requirement already satisfied: typing-extensions>=4.8.0 in /usr/local/lib/python3.10/dist-packages (from torch>=1.10.0->accelerate) (4.11.0)
 Requirement already satisfied: sympy in /usr/local/lib/python3.10/dist-packages (from torch>=1.10.0->accelerate) (1.12)
 Requirement already satisfied: networkx in /usr/local/lib/python3.10/dist-packages (from torch>=1.10.0->accelerate) (3.3)
 Requirement already satisfied: jinja2 in /usr/local/lib/python3.10/dist-packages (from torch>=1.10.0->accelerate) (3.1.3)
 Requirement already satisfied: fsspec in /usr/local/lib/python3.10/dist-packages (from torch>=1.10.0->accelerate) (2023.6.0)
 Requirement already satisfied: nvidia-cuda-nvrtc-cu12==12.1.105 in /usr/local/lib/python3.10/dist-packages (from torch>=1.10.0->accelerate) (12.1.105)
 Requirement already satisfied: nvidia-cuda-runtime-cu12==12.1.105 in /usr/local/lib/python3.10/dist-packages (from torch>=1.10.0->accelerate) (12.1.105)
 Requirement already satisfied: nvidia-cuda-cupti-cu12==12.1.105 in /usr/local/lib/python3.10/dist-packages (from torch>=1.10.0->accelerate) (12.1.105)
 Requirement already satisfied: nvidia-cudnn-cu12==8.9.2.26 in /usr/local/lib/python3.10/dist-packages (from torch>=1.10.0->accelerate) (8.9.2.26)
 Requirement already satisfied: nvidia-cublas-cu12==12.1.3.1 in /usr/local/lib/python3.10/dist-packages (from torch>=1.10.0->accelerate) (12.1.3.1)
 Requirement already satisfied: nvidia-cufft-cu12==11.0.2.54 in /usr/local/lib/python3.10/dist-packages (from torch>=1.10.0->accelerate) (11.0.2.54)
 Requirement already satisfied: nvidia-curand-cu12==10.3.2.106 in /usr/local/lib/python3.10/dist-packages (from torch>=1.10.0->accelerate)

(10.3.2.106)
Requirement already satisfied: nvidia-cusolver-cu12==11.4.5.107 in
/usr/local/lib/python3.10/dist-packages (from torch>=1.10.0->accelerate)
(11.4.5.107)
Requirement already satisfied: nvidia-cuspars-cu12==12.1.0.106 in
/usr/local/lib/python3.10/dist-packages (from torch>=1.10.0->accelerate)
(12.1.0.106)
Requirement already satisfied: nvidia-nccl-cu12==2.19.3 in
/usr/local/lib/python3.10/dist-packages (from torch>=1.10.0->accelerate)
(2.19.3)
Requirement already satisfied: nvidia-nvtx-cu12==12.1.105 in
/usr/local/lib/python3.10/dist-packages (from torch>=1.10.0->accelerate)
(12.1.105)
Requirement already satisfied: triton==2.2.0 in /usr/local/lib/python3.10/dist-
packages (from torch>=1.10.0->accelerate) (2.2.0)
Requirement already satisfied: nvidia-nvjitlink-cu12 in
/usr/local/lib/python3.10/dist-packages (from nvidia-cusolver-
cu12==11.4.5.107->torch>=1.10.0->accelerate) (12.4.127)
Requirement already satisfied: requests in /usr/local/lib/python3.10/dist-
packages (from huggingface-hub->accelerate) (2.31.0)
Requirement already satisfied: tqdm==4.42.1 in /usr/local/lib/python3.10/dist-
packages (from huggingface-hub->accelerate) (4.66.2)
Requirement already satisfied: MarkupSafe>=2.0 in
/usr/local/lib/python3.10/dist-packages (from jinja2->torch>=1.10.0->accelerate)
(2.1.5)
Requirement already satisfied: charset-normalizer<4,>=2 in
/usr/local/lib/python3.10/dist-packages (from requests->huggingface-
hub->accelerate) (3.3.2)
Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.10/dist-
packages (from requests->huggingface-hub->accelerate) (3.6)
Requirement already satisfied: urllib3<3,>=1.21.1 in
/usr/local/lib/python3.10/dist-packages (from requests->huggingface-
hub->accelerate) (2.0.7)
Requirement already satisfied: certifi>=2017.4.17 in
/usr/local/lib/python3.10/dist-packages (from requests->huggingface-
hub->accelerate) (2024.2.2)
Requirement already satisfied: mpmath>=0.19 in /usr/local/lib/python3.10/dist-
packages (from sympy->torch>=1.10.0->accelerate) (1.3.0)

```
[5]: import os
import numpy as np
import datasets
import pandas as pd
from PIL import Image
from pathlib import Path
from tqdm.auto import tqdm
import multiprocessing as mp
```

```

import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split

import torch
import torch.nn as nn
import torch.nn.functional as F
from torchvision import io, transforms
from torch.utils.data import Dataset, DataLoader, random_split

from transformers import Seq2SeqTrainer ,Seq2SeqTrainingArguments
from transformers import VisionEncoderDecoderModel , ViTFeatureExtractor
from transformers import AutoTokenizer , GPT2Config , default_data_collator

if torch.cuda.is_available():

    device = torch.device("cuda")

    print('There are %d GPU(s) available.' % torch.cuda.device_count())

    print('We will use the GPU:', torch.cuda.get_device_name(0))

else:
    print('No GPU available, using the CPU instead.')
    device = torch.device("cpu")

```

There are 1 GPU(s) available.
We will use the GPU: Tesla T4

```

[6]: os.environ["WANDB_DISABLED"] = "true"
class config :
    ENCODER = "google/vit-base-patch16-224"
    DECODER = "gpt2"
    TRAIN_BATCH_SIZE = 8
    VAL_BATCH_SIZE = 8
    VAL_EPOCHS = 1
    LR = 5e-5
    SEED = 42
    MAX_LEN = 128
    SUMMARY_LEN = 20
    WEIGHT_DECAY = 0.01
    MEAN = (0.485, 0.456, 0.406)
    STD = (0.229, 0.224, 0.225)
    TRAIN_PCT = 0.95
    NUM_WORKERS = mp.cpu_count()
    EPOCHS = 3
    IMG_SIZE = (224,224)

```



```

LABEL_MASK = -100
TOP_K = 1000
TOP_P = 0.95

```

```

[7]: def build_inputs_with_special_tokens(self, token_ids_0, token_ids_1=None):
      outputs = [self.bos_token_id] + token_ids_0 + [self.eos_token_id]
      return outputs
AutoTokenizer.build_inputs_with_special_tokens = \
↳ build_inputs_with_special_tokens

```

```

[8]: rouge = datasets.load_metric("rouge")

def compute_metrics(pred):
    labels_ids = pred.label_ids
    pred_ids = pred.predictions

    # all unnecessary tokens are removed
    pred_str = tokenizer.batch_decode(pred_ids, skip_special_tokens=True)
    labels_ids[labels_ids == -100] = tokenizer.pad_token_id
    label_str = tokenizer.batch_decode(labels_ids, skip_special_tokens=True)

    rouge_output = rouge.compute(predictions=pred_str, references=label_str, \
↳ rouge_types=["rouge2"])["rouge2"].mid

    return {
        "rouge2_precision": round(rouge_output.precision, 4),
        "rouge2_recall": round(rouge_output.recall, 4),
        "rouge2_fmeasure": round(rouge_output.fmeasure, 4),
    }

```

```

<ipython-input-8-15d4c46f5e43>:1: FutureWarning: load_metric is deprecated and
will be removed in the next major version of datasets. Use 'evaluate.load'
instead, from the new library Evaluate: https://huggingface.co/docs/evaluate
rouge = datasets.load_metric("rouge")
/usr/local/lib/python3.10/dist-packages/datasets/load.py:756: FutureWarning: The
repository for rouge contains custom code which must be executed to correctly
load the metric. You can inspect the repository content at https://raw.githubusercontent.com/huggingface/datasets/2.18.0/metrics/rouge/rouge.py
You can avoid this message in future by passing the argument
`trust_remote_code=True`.
Passing `trust_remote_code=True` will be mandatory to load this metric from the
next major release of `datasets`.
warnings.warn(
Downloading builder script: 0%|          | 0.00/2.17k [00:00<?, ?B/s]

```

```
[9]: feature_extractor = ViTFeatureExtractor.from_pretrained(config.ENCODER)
tokenizer = AutoTokenizer.from_pretrained(config.DECODER)
tokenizer.pad_token = tokenizer.unk_token
```

```
/usr/local/lib/python3.10/dist-packages/huggingface_hub/utils/_token.py:88:
UserWarning:
The secret `HF_TOKEN` does not exist in your Colab secrets.
To authenticate with the Hugging Face Hub, create a token in your settings tab
(https://huggingface.co/settings/tokens), set it as secret in your Google Colab
and restart your session.
You will be able to reuse this secret in all of your notebooks.
Please note that authentication is recommended but still optional to access
public models or datasets.
```

```
warnings.warn(

preprocessor_config.json: 0%|          | 0.00/160 [00:00<?, ?B/s]

/usr/local/lib/python3.10/dist-
packages/transformers/models/vit/feature_extraction_vit.py:28: FutureWarning:
The class ViTFeatureExtractor is deprecated and will be removed in version 5 of
Transformers. Please use ViTImageProcessor instead.
warnings.warn(

tokenizer_config.json: 0%|          | 0.00/26.0 [00:00<?, ?B/s]
config.json: 0%|          | 0.00/665 [00:00<?, ?B/s]
vocab.json: 0%|          | 0.00/1.04M [00:00<?, ?B/s]
merges.txt: 0%|          | 0.00/456k [00:00<?, ?B/s]
tokenizer.json: 0%|          | 0.00/1.36M [00:00<?, ?B/s]
```

```
[11]: from google.colab import drive
drive.mount('/content/drive')
```

Mounted at /content/drive

```
[12]: import zipfile

# Specify the path to the uploaded zip file
zip_file_path = 'drive/MyDrive/custom_captions_dataset.zip'

# Specify the directory where you want to extract the contents
extracti_dir = '/content/custom_captions_dataset'

# Create the extract directory if it doesn't exist
os.makedirs(extracti_dir, exist_ok=True)

# Extract the zip file
with zipfile.ZipFile(zip_file_path, 'r') as zip_ref:
```

```
zip_ref.extractall(extracti_dir)
```

```
[13]: import torchvision.transforms as transforms

transforms = transforms.Compose(
    [
        transforms.Resize(config.IMG_SIZE),
        transforms.ToTensor(),
        transforms.Normalize(
            mean=0.5,
            std=0.5
        )
    ]
)

df1= pd.read_csv("/content/custom_captions_dataset/custom_captions_dataset/
    ↪train.csv")
train_df = df1.iloc[:, 1:3]

df2= pd.read_csv("/content/custom_captions_dataset/custom_captions_dataset/val.
    ↪csv")
val_df = df2.iloc[:, 1:3]
```

```
[14]: class ImgDataset(Dataset):
    def __init__(self, df, root_dir, tokenizer, feature_extractor, transform =
    ↪None):
        self.df = df
        self.transform = transform
        self.root_dir = root_dir
        self.tokenizer= tokenizer
        self.feature_extractor = feature_extractor
        self.max_length = 300
    def __len__(self,):
        return len(self.df)
    def __getitem__(self,idx):
        caption = self.df.caption.iloc[idx]
        image = self.df.filename.iloc[idx]
        img_path = os.path.join(self.root_dir,image)
        img = Image.open(img_path).convert("RGB")

        if self.transform is not None:
            img= self.transform(img)

        img = (img+1)/2
        pixel_values = self.feature_extractor(img, return_tensors="pt").
    ↪pixel_values
        captions = self.tokenizer(caption,
```

```

padding='max_length',
max_length=self.max_length).input_ids
captions = [caption if caption != self.tokenizer.pad_token_id else -100
↳for caption in captions]
encoding = {"pixel_values": pixel_values.squeeze(), "labels": torch.
↳tensor(captions)}
return encoding

```

```

[15]: train_dataset = ImgDataset(train_df, root_dir = "/content/
↳custom_captions_dataset/custom_captions_dataset/
↳train",tokenizer=tokenizer,feature_extractor = feature_extractor ,transform_
↳= transforms)
val_dataset = ImgDataset(val_df , root_dir = "/content/custom_captions_dataset/
↳custom_captions_dataset/val",tokenizer=tokenizer,feature_extractor =
↳feature_extractor , transform = transforms)

```

```

[16]: model = VisionEncoderDecoderModel.from_encoder_decoder_pretrained(config.
↳ENCODER, config.DECODER)

```

```
config.json: 0%|          | 0.00/69.7k [00:00<?, ?B/s]
```

```
model.safetensors: 0%|          | 0.00/346M [00:00<?, ?B/s]
```

Some weights of ViTModel were not initialized from the model checkpoint at google/vit-base-patch16-224 and are newly initialized: ['vit.pooler.dense.bias', 'vit.pooler.dense.weight']

You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference.

```
model.safetensors: 0%|          | 0.00/548M [00:00<?, ?B/s]
```

Some weights of GPT2LMHeadModel were not initialized from the model checkpoint at gpt2 and are newly initialized: ['h.0.crossattention.c_attn.bias', 'h.0.crossattention.c_attn.weight', 'h.0.crossattention.c_proj.bias', 'h.0.crossattention.c_proj.weight', 'h.0.crossattention.q_attn.bias', 'h.0.crossattention.q_attn.weight', 'h.0.ln_cross_attn.bias', 'h.0.ln_cross_attn.weight', 'h.1.crossattention.c_attn.bias', 'h.1.crossattention.c_attn.weight', 'h.1.crossattention.c_proj.bias', 'h.1.crossattention.c_proj.weight', 'h.1.crossattention.q_attn.bias', 'h.1.crossattention.q_attn.weight', 'h.1.ln_cross_attn.bias', 'h.1.ln_cross_attn.weight', 'h.10.crossattention.c_attn.bias', 'h.10.crossattention.c_attn.weight', 'h.10.crossattention.c_proj.bias', 'h.10.crossattention.c_proj.weight', 'h.10.crossattention.q_attn.bias', 'h.10.crossattention.q_attn.weight', 'h.10.ln_cross_attn.bias', 'h.10.ln_cross_attn.weight', 'h.11.crossattention.c_attn.bias', 'h.11.crossattention.c_attn.weight', 'h.11.crossattention.c_proj.bias', 'h.11.crossattention.c_proj.weight', 'h.11.crossattention.q_attn.bias', 'h.11.crossattention.q_attn.weight', 'h.11.ln_cross_attn.bias', 'h.11.ln_cross_attn.weight', 'h.2.crossattention.c_attn.bias',

```

'h.2.crossattention.c_attn.weight', 'h.2.crossattention.c_proj.bias',
'h.2.crossattention.c_proj.weight', 'h.2.crossattention.q_attn.bias',
'h.2.crossattention.q_attn.weight', 'h.2.ln_cross_attn.bias',
'h.2.ln_cross_attn.weight', 'h.3.crossattention.c_attn.bias',
'h.3.crossattention.c_attn.weight', 'h.3.crossattention.c_proj.bias',
'h.3.crossattention.c_proj.weight', 'h.3.crossattention.q_attn.bias',
'h.3.crossattention.q_attn.weight', 'h.3.ln_cross_attn.bias',
'h.3.ln_cross_attn.weight', 'h.4.crossattention.c_attn.bias',
'h.4.crossattention.c_attn.weight', 'h.4.crossattention.c_proj.bias',
'h.4.crossattention.c_proj.weight', 'h.4.crossattention.q_attn.bias',
'h.4.crossattention.q_attn.weight', 'h.4.ln_cross_attn.bias',
'h.4.ln_cross_attn.weight', 'h.5.crossattention.c_attn.bias',
'h.5.crossattention.c_attn.weight', 'h.5.crossattention.c_proj.bias',
'h.5.crossattention.c_proj.weight', 'h.5.crossattention.q_attn.bias',
'h.5.crossattention.q_attn.weight', 'h.5.ln_cross_attn.bias',
'h.5.ln_cross_attn.weight', 'h.6.crossattention.c_attn.bias',
'h.6.crossattention.c_attn.weight', 'h.6.crossattention.c_proj.bias',
'h.6.crossattention.c_proj.weight', 'h.6.crossattention.q_attn.bias',
'h.6.crossattention.q_attn.weight', 'h.6.ln_cross_attn.bias',
'h.6.ln_cross_attn.weight', 'h.7.crossattention.c_attn.bias',
'h.7.crossattention.c_attn.weight', 'h.7.crossattention.c_proj.bias',
'h.7.crossattention.c_proj.weight', 'h.7.crossattention.q_attn.bias',
'h.7.crossattention.q_attn.weight', 'h.7.ln_cross_attn.bias',
'h.7.ln_cross_attn.weight', 'h.8.crossattention.c_attn.bias',
'h.8.crossattention.c_attn.weight', 'h.8.crossattention.c_proj.bias',
'h.8.crossattention.c_proj.weight', 'h.8.crossattention.q_attn.bias',
'h.8.crossattention.q_attn.weight', 'h.8.ln_cross_attn.bias',
'h.8.ln_cross_attn.weight', 'h.9.crossattention.c_attn.bias',
'h.9.crossattention.c_attn.weight', 'h.9.crossattention.c_proj.bias',
'h.9.crossattention.c_proj.weight', 'h.9.crossattention.q_attn.bias',
'h.9.crossattention.q_attn.weight', 'h.9.ln_cross_attn.bias',
'h.9.ln_cross_attn.weight']

```

You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference.

```
generation_config.json: 0%|          | 0.00/124 [00:00<?, ?B/s]
```

```

[17]: model.config.decoder_start_token_id = tokenizer.cls_token_id
model.config.pad_token_id = tokenizer.pad_token_id
# make sure vocab size is set correctly
model.config.vocab_size = model.config.decoder.vocab_size
# set beam search parameters
model.config.eos_token_id = tokenizer.sep_token_id
model.config.decoder_start_token_id = tokenizer.bos_token_id
model.config.max_length = 128
model.config.early_stopping = True
model.config.no_repeat_ngram_size = 3
model.config.length_penalty = 2.0

```

```
model.config.num_beams = 4
```

```
[18]: from transformers import Seq2SeqTrainingArguments

training_args = Seq2SeqTrainingArguments(
    output_dir='VIT_large_gpt2',
    per_device_train_batch_size=config.TRAIN_BATCH_SIZE,
    per_device_eval_batch_size=config.VAL_BATCH_SIZE,
    predict_with_generate=True,
    evaluation_strategy="epoch",
    do_train=True,
    do_eval=True,
    logging_steps=1024,
    save_steps=2048,
    warmup_steps=1024,
    learning_rate = 5e-5,
    #max_steps=1500, # delete for full training
    num_train_epochs = config.EPOCHS, #TRAIN_EPOCHS
    overwrite_output_dir=True,
    save_total_limit=1,
)
```

Using the `WANDB_DISABLED` environment variable is deprecated and will be removed in v5. Use the `--report_to` flag to control the integrations used for logging result (for instance `--report_to none`).

```
[19]: trainer=Seq2SeqTrainer(
    tokenizer=feature_extractor,
    model=model,
    args=training_args,
    compute_metrics=compute_metrics,
    train_dataset=train_dataset,
    eval_dataset=val_dataset,
    data_collator=default_data_collator,
)
trainer.train()
```

```
/usr/local/lib/python3.10/dist-packages/accelerate/accelerator.py:436:
FutureWarning: Passing the following arguments to `Accelerator` is deprecated
and will be removed in version 1.0 of Accelerate: dict_keys(['dispatch_batches',
'split_batches', 'even_batches', 'use_seedable_sampler']). Please pass an
`accelerate.DataLoaderConfiguration` instead:
dataloader_config = DataLoaderConfiguration(dispatch_batches=None,
split_batches=False, even_batches=True, use_seedable_sampler=True)
warnings.warn(
It looks like you are trying to rescale already rescaled images. If the input
```

images have pixel values between 0 and 1, set ``do_rescale=False`` to avoid rescaling them again.

<IPython.core.display.HTML object>

```
/usr/local/lib/python3.10/dist-packages/transformers/generation/utils.py:1339:
UserWarning: You have modified the pretrained model configuration to control
generation. This is a deprecated strategy to control generation and will be
removed soon, in a future version. Please use and modify the model generation
configuration (see
https://huggingface.co/docs/transformers/generation_strategies#default-text-
generation-configuration )
```

```
warnings.warn(
Some non-default generation parameters are set in the model config. These should
go into a GenerationConfig file
(https://huggingface.co/docs/transformers/generation_strategies#save-a-custom-
decoding-strategy-with-your-model) instead. This warning will be raised to an
exception in v4.41.
```

```
Non-default generation parameters: {'max_length': 128, 'early_stopping': True,
'num_beams': 4, 'length_penalty': 2.0, 'no_repeat_ngram_size': 3}
```

```
Removed shared tensor {'decoder.lm_head.weight'} while saving. This should be
OK, but check by verifying that you don't receive any warning while reloading
```

```
[19]: TrainOutput(global_step=2145, training_loss=2.514048052167559,
metrics={'train_runtime': 5856.3062, 'train_samples_per_second': 2.928,
'train_steps_per_second': 0.366, 'total_flos': 3.0940534132545946e+18,
'train_loss': 2.514048052167559, 'epoch': 3.0})
```

```
[20]: trainer.save_model('VIT_large_gpt2')
```

```
Some non-default generation parameters are set in the model config. These should
go into a GenerationConfig file
(https://huggingface.co/docs/transformers/generation_strategies#save-a-custom-
decoding-strategy-with-your-model) instead. This warning will be raised to an
exception in v4.41.
```

```
Non-default generation parameters: {'max_length': 128, 'early_stopping': True,
'num_beams': 4, 'length_penalty': 2.0, 'no_repeat_ngram_size': 3}
```

```
[21]: img_ = Image.open("/content/custom_captions_dataset/custom_captions_dataset/
↳test/test_25.jpg").convert("RGB")
img_
```

```
[21]:
```



```
[22]: generated_caption = tokenizer.decode(model.generate(feature_extractor(img_,
    ↪return_tensors="pt").pixel_values.to("cuda"))[0])
print('\033[96m' + generated_caption[:142] + '\033[0m')
```

<|endoftext|>A man is standing on a tennis court. He is wearing a white shirt and blue shorts. The man is holding a racket in his right hand.

```
[23]: df3 = pd.read_csv("/content/custom_captions_dataset/custom_captions_dataset/
    ↪test.csv")
test_df=df3.iloc[:, 1:3]

# Define the dataset using ImgDataset class
test_dataset = ImgDataset(test_df, root_dir="/content/custom_captions_dataset/
    ↪custom_captions_dataset/test", tokenizer=tokenizer,
    ↪feature_extractor=feature_extractor, transform=transforms)

# Define a function to generate captions for test images
def generate_captions(dataset):
    captions = []
    for idx in range(len(dataset)):
        inputs = dataset[idx]
        img_tensor = inputs["pixel_values"].unsqueeze(0).to(device)
        generated_caption = tokenizer.decode(model.generate(img_tensor)[0])
```



```

        captions.append(generated_caption)
    return captions

# Generate captions for test dataset
generated_captions = generate_captions(test_dataset)

# Assuming you have ground truth captions for evaluation
ground_truth_captions = test_df["caption"].tolist()

```

```

[43]: !git clone https://github.com/vrama91/cider.git
      %cd cider
      !python setup.py install

```

```

Cloning into 'cider'...
remote: Enumerating objects: 218, done.
remote: Total 218 (delta 0), reused 0 (delta 0), pack-reused 218
Receiving objects: 100% (218/218), 31.08 MiB | 15.25 MiB/s, done.
Resolving deltas: 100% (118/118), done.
/content/cider
python3: can't open file '/content/cider/setup.py': [Errno 2] No such file or
directory

```

```

[3]: from cider.pyciderevalcap.cider.cider import Cider

# Assuming df_mode represents the mode of your dataframe, such as 'train',
    ↪ 'val', or 'test'
df_mode = "test"

ground_truth_data = [{'image_id': i, 'caption': caption} for i, caption in
    ↪ enumerate(ground_truth_captions)]

# Format generated captions
generated_data = [{'image_id': i, 'caption': caption} for i, caption in
    ↪ enumerate(generated_captions_list)]

# Print the first few elements of ground_truth_data and generated_data
print("First few elements of ground_truth_data:", ground_truth_data[:5])
print("First few elements of generated_data:", generated_data[:5])

# Check the type of the first caption in ground_truth_data and generated_data
print("Type of the first caption in ground_truth_data:",
    ↪ type(ground_truth_data[0]['caption']))
print("Type of the first caption in generated_data:",
    ↪ type(generated_data[0]['caption']))

# Format ground truth captions as lists of strings
for item in ground_truth_data:

```

```

    item['caption'] = [item['caption']]

# Format generated captions as lists of strings
for item in generated_data:
    item['caption'] = [item['caption']]

# Initialize the Cider with the appropriate arguments
cider = Cider(df_mode)

# Now you can compute the CIDEr scores
cider_scores = cider.compute_score(ground_truth_data, generated_data)
cider_score = sum(cider_scores) / len(cider_scores)
print("CIDEr Score:", cider_score)

```

CIDEr Score: 0.24473210931828393

[33]: !pip install rouge

```

Collecting rouge
  Downloading rouge-1.0.1-py3-none-any.whl (13 kB)
Requirement already satisfied: six in /usr/local/lib/python3.10/dist-packages
(from rouge) (1.16.0)
Installing collected packages: rouge
Successfully installed rouge-1.0.1

```

```

[34]: from rouge import Rouge

rouge = Rouge()
rouge_scores = rouge.get_scores(generated_captions, ground_truth_captions,
    ↪ avg=True)
rouge_l_score = rouge_scores["rouge-1"]["f"]
print("ROUGE-L Score:", rouge_l_score)

```

ROUGE-L Score: 0.2873843743802736

[36]: !pip install SpicePy

```

Collecting SpicePy
  Downloading spicepy-1.0.5-py3-none-any.whl (17 kB)
Requirement already satisfied: numpy in /usr/local/lib/python3.10/dist-packages
(from SpicePy) (1.25.2)
Requirement already satisfied: scipy in /usr/local/lib/python3.10/dist-packages
(from SpicePy) (1.11.4)
Requirement already satisfied: matplotlib in /usr/local/lib/python3.10/dist-
packages (from SpicePy) (3.7.1)
Requirement already satisfied: contourpy>=1.0.1 in
/usr/local/lib/python3.10/dist-packages (from matplotlib->SpicePy) (1.2.1)
Requirement already satisfied: cycycler>=0.10 in /usr/local/lib/python3.10/dist-

```

```
packages (from matplotlib->SpicePy) (0.12.1)
Requirement already satisfied: fonttools>=4.22.0 in
/usr/local/lib/python3.10/dist-packages (from matplotlib->SpicePy) (4.51.0)
Requirement already satisfied: kiwisolver>=1.0.1 in
/usr/local/lib/python3.10/dist-packages (from matplotlib->SpicePy) (1.4.5)
Requirement already satisfied: packaging>=20.0 in
/usr/local/lib/python3.10/dist-packages (from matplotlib->SpicePy) (24.0)
Requirement already satisfied: pillow>=6.2.0 in /usr/local/lib/python3.10/dist-
packages (from matplotlib->SpicePy) (9.4.0)
Requirement already satisfied: pyparsing>=2.3.1 in
/usr/local/lib/python3.10/dist-packages (from matplotlib->SpicePy) (3.1.2)
Requirement already satisfied: python-dateutil>=2.7 in
/usr/local/lib/python3.10/dist-packages (from matplotlib->SpicePy) (2.8.2)
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.10/dist-
packages (from python-dateutil>=2.7->matplotlib->SpicePy) (1.16.0)
Installing collected packages: SpicePy
Successfully installed SpicePy-1.0.5
```

```
[2]: import spicepy
      # Initialize the SPICE scorer

      spice = spicepy.Spice()

      # Compute SPICE scores
      spice_scores = spice.compute_score(ground_truth_captions, generated_captions)

      # Extract the overall SPICE score
      spice_score = spice_scores[0]["All"]["f"]

      print("SPICE Score:", spice_score)
```

SPICE Score: 0.2913853228547193

```
[ ]:
```